

ABDUL QADIR

Full Stack AI Engineer | AI Backend Engineer | Generative AI Engineer | NLP & Computer Vision
+92 321 331 3138 | abdulqadirmagssii@gmail.com | [LinkedIn](#) | [GitHub](#) | Islamabad, Pakistan

PROFESSIONAL SUMMARY

Results-driven AI Engineer with 3+ years of experience building Generative AI systems, LLM applications, and computer vision pipelines. Skilled in RAG architectures, LangChain/LangGraph workflows, and model fine-tuning using PyTorch/TensorFlow. Experienced in building scalable AI system architectures. Currently developing agentic AI products serving 10,000+ users, delivering end-to-end solutions from data processing to production.

TECHNICAL SKILLS

Languages: Python, C++, Java, HTML, CSS

Generative AI & LLMs: GPT-4, LLaMA, Mistral, Gemma, Phi, Claude, Prompt Engineering, LLM Fine-Tuning, RLHF, PEFT, LoRA

RAG & Vector Databases: LangChain, LlamaIndex, RAG Pipelines, Pinecone, FAISS, Chroma, Weaviate

Agentic AI: LangGraph, CrewAI, Multi-Agent Systems, AI Workflow Automation (n8n)

ML/DL Frameworks: PyTorch, TensorFlow, Scikit-learn, Keras, RNN, LSTM, CNN, BERT, ResNet, YOLO v8/v9

MLOps & DevOps: Docker, Docker Hub, Azure ML, LangSmith, MLflow, CI/CD Pipelines, Roboflow

Frameworks & APIs: FastAPI, Flask, Django, Streamlit, Gradio

WORK EXPERIENCE

Prompt & Generative AI Engineer (Full-Time) | Sideline Technologies (PVT) LTD | Jul 2025 – Present | Onsite

- Engineered large-scale, production-ready AI products serving 10,000+ active users by applying advanced prompt engineering, vibe coding workflows, and rapid iterative development cycles.
- Built end-to-end RAG systems using LangChain, FAISS, and Chroma to enable semantic search and intelligent document Q&A across large proprietary knowledge bases.
- Designed multi-step agentic workflows with LangGraph and external API integrations to automate complex user-facing tasks, reducing manual processing time significantly.
- Optimized LLM prompt strategies (chain-of-thought, few-shot, system prompt tuning) and monitored production performance via LangSmith to maintain quality at scale.

AI Engineer (Full-Time) | ErlyStage Studio | Sep 2024 – Jun 2025 | Remote

- Engineered production-ready RAG pipelines integrating LangChain, FAISS, and Chroma vector databases, reducing document retrieval latency by approximately 35%.
- Built multi-agent AI workflows using LangGraph and CrewAI, automating NLP tasks including sentiment analysis, entity extraction, and classification across 10+ business use cases.
- Fine-tuned open-source LLMs (LLaMA, Mistral) using PEFT/LoRA techniques, improving task accuracy by up to 25% over baseline models.
- Containerized and deployed AI microservices with Docker, establishing CI/CD-compatible model serving pipelines.

Generative AI Engineer (Full-Time) | Xflow Research Inc. | Aug 2023 – Jun 2024 | Hybrid

- Led applied research on LLM fine-tuning, RAG architectures, and agentic AI systems, contributing to 3 production deployments for enterprise clients.
- Developed prompt engineering strategies for GPT-4 and open-source LLMs using chain-of-thought and few-shot techniques, measurably improving output quality.
- Built end-to-end RAG systems with Pinecone and LlamaIndex enabling semantic search over large proprietary document corpora.
- Collaborated on multimodal AI prototypes combining vision and language models for automated document understanding pipelines.

Computer Vision Engineer (Internship) | National Center for Robotics & Automation | May 2023 – Jul 2023 | Pakistan

- Implemented real-time object detection and instance segmentation pipelines using YOLO v8, MobileNetV2, and ResNet-based CNN architectures.
- Deployed trained vision models on Raspberry Pi and NVIDIA Jetson Nano, achieving real-time inference at 25+ FPS for robotics applications.

KEY PROJECTS

PatchPilot — AI Agent Tool Orchestration Layer

- Built an orchestration layer that lets LLM agents use real tools safely: typed tool registry with per-tool scopes, secrets isolation, and rate limits, plus a LangGraph planner with retries, fallbacks, and human-approval checkpoints.
- Stack: Python, FastAPI, LangGraph, LangSmith, Docker, Redis | Outcome: 40+ tools in registry, <300ms routing overhead per step, 99.5% of runs completed without manual retry.

CodeGuard — AI Security & Code Quality Scanner

- Built a scanner combining rule-based detection of secrets, vulnerable dependencies, and injection flaws (Python, JS/TS) with LLM triage using code-context retrieval, posting ranked findings and fix suggestions as PR comments.
- Stack: Python, FastAPI, LangChain, GPT-4, FAISS, GitHub Apps, Docker | Outcome: ~70% less alert noise vs. raw scanner output, 2-min median scan per PR, 1-click patch suggestions.

CTA — AI Trading Journal

- Built an AI trading journal that auto-imports broker trades, tags setups and emotions, supports natural-language queries over trade history, and generates weekly AI coach reports grounded in the user's own trades.
- Stack: Python, FastAPI, LangChain, Chroma, PostgreSQL, React | Outcome: 12k trades journaled in beta, 5-sec trade logging, weekly AI performance reviews.

Multi-Agent Medical AI System

- Built a multi-agent system using LangGraph and CrewAI with specialized agents for diagnosis, triage, and report generation powered by fine-tuned LLMs.
- Stack: Python, LangGraph, CrewAI, LangChain, GPT-4, FAISS | Outcome: Reduced manual clinical review time by ~40% in pilot testing.

Enterprise RAG Document Chat Platform

- Designed a scalable RAG pipeline with multi-format ingestion, semantic chunking, re-ranking, and a conversational UI layer.
- Stack: LangChain, LlamaIndex, Chroma, Gemma 2B, FastAPI, Streamlit | Outcome: Sub-second query responses over 10,000+ document collections.

Real-Time Computer Vision Tracking System

- Implemented multi-object tracking with DeepSORT and YOLO v8 for instance segmentation and speed estimation, deployed on edge hardware.
- Stack: Python, OpenCV, YOLO v8, DeepSORT, NVIDIA Jetson Nano | Outcome: 30 FPS with >90% tracking accuracy.

EDUCATION

Bachelor of Engineering — Computer Systems Engineering

Mehran University of Engineering & Technology, Jamshoro, Pakistan | Dec 2020 – Dec 2024

CERTIFICATIONS

- Intermediate Machine Learning — Kaggle
- Large Language Models: Concepts & Applications — DataCamp
- Recurrent Neural Networks (RNN) for Language Modeling — DataCamp
- Feature Engineering for NLP — DataCamp
- Artificial Intelligence — NAVTTC